

Christopher Miao  
Sulekha Abas

A.

This is our nested loop join showing that and the query performance is time of 5.366 ms.

### NESTED LOOP JOIN: Q1

```
OUTPUT  DEBUG CONSOLE  PROBLEMS  PORTS  TERMINAL
Nested Loop join

psql:/xtra/abas0013/csci5708/Lab2/dataset/airport.sql:2290607: ERROR:  multiple
primary keys for table "tickets" are not allowed
ALTER DATABASE
ALTER DATABASE
demo=# \timing
Timing is on.
demo=# select * from Seats join Aircrafts on Seats.aircraft_code = Aircrafts.ai
rcraft_code limit 5;
 aircraft_code | seat_no | fare_conditions | aircraft_code | model | r
ange
-----+-----+-----+-----+-----+-----
773            | 11A     | Comfort         | 773           | Boeing 777-300 | 1
1100
773            | 11C     | Comfort         | 773           | Boeing 777-300 | 1
1100
773            | 11D     | Comfort         | 773           | Boeing 777-300 | 1
1100
773            | 11E     | Comfort         | 773           | Boeing 777-300 | 1
1100
773            | 11F     | Comfort         | 773           | Boeing 777-300 | 1
1100
(5 rows)

Time: 5.366 ms
demo=#
```

This is our merge join with query performance of 575.627 ms.

### MERGE JOIN: Q2

```
File Edit View Search Terminal Help
LOG: autovacuum launcher started
ERROR: invalid input syntax for type json
DETAIL: Character with value 0x11 must be escaped.
CONTEXT: JSON data, line 1: {"en": "Boeing 777-300"}
COPY airports data, line 1, column model:
STATEMENT: COPY airports data (airport_code, model)
ERROR: invalid input syntax for type json
DETAIL: Character with value 0x1c must be escaped.
CONTEXT: JSON data, line 1: {"en": "Mirny Airport"}
COPY airports data, line 2, column airport_n
STATEMENT: COPY airports data (airport_code, airport_n)
ERROR: table name "bookings" specified more than on
STATEMENT: SELECT * FROM Bookings JOIN Bookings ON
;
nest loop init
ERROR: syntax error at or near "SELECT" at character
STATEMENT: SELECT * FROM Ticket_flights JOIN Boarding
LIMIT 3
SELECT * FROM Ticket_flights JOIN Boarding_o
IT 3;
nest loop init
merge join init
merge join init

demo=# SELECT * from (SELECT * FROM Bookings Order by book_ref ASC) as b JOIN (SELECT * FROM Tickets Ord
er by book_ref ASC) as t on b.book_ref = t.book_ref LIMIT 3;
Time: 598.588 ms
demo=# \timing
Timing is off.
demo=# \timing
Timing is on.
demo=# SELECT * from (SELECT * FROM Bookings Order by book_ref ASC) as b JOIN (SELECT * FROM Tickets Ord
er by book_ref ASC) as t on b.book_ref = t.book_ref LIMIT 3;
 book_ref | book_date | total_amount | ticket_no | book_ref | passenger_id | passenger_
name | contact_data
-----+-----+-----+-----+-----+-----+-----
00000F | 2017-07-04 19:12:00-05 | 265700.00 | 0005435838975 | 00000F | 1708 262537 | ANNA ANTON
OVA | {"email": "annaantonova-19021973@postgrespro.ru", "phone": "+70938049942"}
000012 | 2017-07-14 01:02:00-05 | 37900.00 | 0005432527326 | 000012 | 9091 269355 | TAMARA ZAY
CEVA | {"email": "tamarazayceva-1971@postgrespro.ru", "phone": "+70749401734"}
000068 | 2017-08-15 06:27:00-05 | 18100.00 | 0005432293273 | 000068 | 5895 674437 | TATYANA PE
TROVA | {"email": "t_petrova1970@postgrespro.ru", "phone": "+70886117503"}
(3 rows)

Time: 575.627 ms
demo=#
```

This is our hash join with query performance of 1.800 ms.

### HASH JOIN: Q3

```
File Edit View Search Terminal Help
DETAIL: Character with value 0x1c must be
CONTEXT: JSON data, line 1: {"en": "Mirny
COPY airports data, line 2, column
STATEMENT: COPY airports_data (airport_cod
ERROR: table name "bookings" specified mor
STATEMENT: SELECT * FROM Bookings JOIN Bo
;
nest loop init
ERROR: syntax error at or near "SELECT" at
STATEMENT: SELECT * FROM Ticket_flights JO
LIMIT 3
SELECT * FROM Ticket_flights JOIN B
IT 3;
nest loop init
merge join init
merge join init
merge join init
nest loop init
nest loop init
hash join init
hash join init
hash join init
hash join init

Time: 0.733 ms
demo=# SELECT * FROM (select * from Bookings limit 1000) as b join (select * from Tickets limit 1000) as
t on b.book_ref = t.book_ref;
 book_ref | book_date | total_amount | ticket_no | book_ref | passenger_id | passenger
name | contact_data
-----+-----+-----+-----+-----+-----+-----
00BBF4 | 2017-07-23 23:27:00-05 | 21400.00 | 0005432006781 | 00BBF4 | 0457 897678 | YURIY KULI
KOV | {"email": "kulikov.y-1979@postgrespro.ru", "phone": "+70902684864"}
00BBF4 | 2017-07-23 23:27:00-05 | 21400.00 | 0005432006782 | 00BBF4 | 4813 271415 | GALINA KUZ
NECOVA | {"phone": "+70741152329"}
005696 | 2017-07-14 12:23:00-05 | 14000.00 | 0005432020118 | 005696 | 8783 431871 | IRINA VASI
LEVA | {"phone": "+70840261333"}
005696 | 2017-07-14 12:23:00-05 | 14000.00 | 0005432020119 | 005696 | 9790 140874 | OLEG SCHER
BAKOV | {"phone": "+70766374027"}
(4 rows)

Time: 1.800 ms
demo=#
```

B.

Below is Query 1, which originally resulted in a Nested Loop join. When run using a nested loop join, it took 0.590 ms. This nested loop join query did not work when we forced it to use a merge join. We tried several other queries that were originally shown to be nested loop joins, but when we tried making those other queries use a merge join, they all would not run. We also tried it on two different machines and faced the same issues. Therefore, we do not know what the exact runtime of an originally nested loop join query would be if it were forced to use a merge join instead. However, when our original nested loop join query was forced to run as a hash merge, it worked fine and had a run time of 0.692 ms.

Nest

```
demo=# select * from Seats join Aircrafts on Seats.aircraft_code = Aircrafts.aircraft_code limit 5;
```

aircraft_code	seat_no	fare_conditions	aircraft_code	model	range
773	11A	Comfort	773		11100
773	11C	Comfort	773		11100
773	11D	Comfort	773		11100
773	11E	Comfort	773		11100
773	11F	Comfort	773		11100

(5 rows)

Time: 0.590 ms  
demo=#

Hash

```
demo=# select * from Seats join Aircrafts on Seats.aircraft_code = Aircrafts.aircraft_code limit 5;
```

aircraft_code	seat_no	fare_conditions	aircraft_code	model	range
319	2A	Business	319		6700
319	2C	Business	319		6700
319	2D	Business	319		6700
319	2F	Business	319		6700
319	3A	Business	319		6700

(5 rows)

Time: 0.692 ms  
demo=#

Below is Query 2, which originally resulted in a Merge join. The time it took for Q2 to run using a nested loop join was 1062.6 ms. When Q2 used a merge join, it took 640.1 ms, and when it used a Hash 1024.3 ms.

demo=# select \* from (select \* from Bookings order by book\_ref asc) as b join (select \* from Tickets order by book\_ref asc) as t on b.book\_ref = t.book\_ref limit 3;

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
000012	2017-07-14 01:02:00-05	37900.00	0005432527326	000012	9091 269355	TAMARA ZAYCEVA	{"email": "tamarazayceva-1971@postgrespro.ru", "phone": "+70749401734"}
000068	2017-08-15 06:27:00-05	18100.00	0005432293273	000068	5895 674437	TATYANA PETROVA	{"email": "t_petrova1970@postgrespro.ru", "phone": "+70886117503"}

(3 rows)

Time: 1062.593 ms

demo=# select \* from (select \* from Bookings order by book\_ref asc) as b join (select \* from Tickets order by book\_ref asc) as t on b.book\_ref = t.book\_ref limit 3;

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
00000F	2017-07-04 19:12:00-05	265700.00	0005435838975	00000F	1708 262537	ANNA ANTONOVA	{"email": "annaantonova-19021973@postgrespro.ru", "phone": "+70938049942"}
000012	2017-07-14 01:02:00-05	37900.00	0005432527326	000012	9091 269355	TAMARA ZAYCEVA	{"email": "tamarazayceva-1971@postgrespro.ru", "phone": "+70749401734"}
000068	2017-08-15 06:27:00-05	18100.00	0005432293273	000068	5895 674437	TATYANA PETROVA	{"email": "t_petrova1970@postgrespro.ru", "phone": "+70886117503"}

(3 rows)

Time: 640.145 ms

demo=# select \* from (select \* from Bookings order by book\_ref asc) as b join (select \* from Tickets order by book\_ref asc) as t on b.book\_ref = t.book\_ref limit 3;

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
000068	2017-08-15 06:27:00-05	18100.00	0005432293273	000068	5895 674437	TATYANA PETROVA	{"email": "t_petrova1970@postgrespro.ru", "phone": "+70886117503"}
0006C3	2017-08-02 03:42:00-05	59400.00	0005432329491	0006C3	4486 459735	YULIYA TROFIMOVA	{"email": "yuliyatrofimova-011988@postgrespro.ru", "phone": "+70850747734"}
0006C3	2017-08-02 03:42:00-05	59400.00	0005432329492	0006C3	1691 126633	FARIT ZAKHAROV	{"email": "zakharov-farit27091972@postgrespro.ru", "phone": "+70381259324"}

(3 rows)

Time: 1024.300 ms

```
nest loop init  
merge join init  
hash join init
```

Below is Query 3, which originally resulted in a Hash join. The time it took for Q3 to run as a nested loop join was 69.150 ms, with a merge join it took 1.420 ms, and with a hash join it took 0.464 ms.

Time: 0.124 ms  
demo=# select \* from (select \* from Bookings limit 1000) as b join (select \* from Tickets limit 1000) as t on b.book\_ref = t.book\_ref;  
ERROR: syntax error at or near "from"  
LINE 1: ... \* from Bookings limit 1000) as b join (select \* from Ticke...

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
005696	2017-07-14 12:23:00-05	14000.00	0005432020118	005696	8783 431871	IRINA VASILEVA	{"phone": "+70840261333"}
005696	2017-07-14 12:23:00-05	14000.00	0005432020119	005696	9790 140874	OLEG SCHERBAKOV	{"phone": "+70766374027"}
008BF4	2017-07-23 23:27:00-05	21400.00	0005432006781	008BF4	0457 897678	YURIY KULIKOV	{"email": "kulikov.y-1979@postgrespro.ru", "phone": "+70902684864"}
008BF4	2017-07-23 23:27:00-05	21400.00	0005432006782	008BF4	4813 271415	GALINA KUZNECOVA	{"phone": "+70741152329"}

(4 rows)

Time: 69.150 ms  
demo=# select \* from (select \* from Bookings limit 1000) as b join (select \* from Tickets limit 1000) as t on b.book\_ref = t.book\_ref;  
Time: 1.420 ms  
demo=# select \* from (select \* from Bookings limit 1000) as b join (select \* from Tickets limit 1000) as t on b.book\_ref = t.book\_ref;  
Time: 0.464 ms

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
005696	2017-07-14 12:23:00-05	14000.00	0005432020118	005696	8783 431871	IRINA VASILEVA	{"phone": "+70840261333"}
005696	2017-07-14 12:23:00-05	14000.00	0005432020119	005696	9790 140874	OLEG SCHERBAKOV	{"phone": "+70766374027"}
000367	2017-08-05 16:05:00-05	6700.00	0005432020751	000367	9176 505722	NIKITA NAUMOV	{"email": "naumov_nikita_061974@postgrespro.ru", "phone": "+70517860377"}

(3 rows)

book_ref	book_date	total_amount	ticket_no	book_ref	passenger_id	passenger_name	contact_data
001269	2017-07-05 18:07:00-05	30000.00	0005432034629	001269	5157 134525	NIKOLAY CHERNOV	{"email": "chernov-n1968@postgrespro.ru", "phone": "+70359486833"}
001269	2017-07-05 18:07:00-05	30000.00	0005432034628	001269	2667 544137	MARIYA NIKOLAEVA	{"phone": "+70508305342"}
001269	2017-07-05 18:07:00-05	30000.00	0005432034627	001269	2853 922120	BORIS ORLOV	{"email": "orlov.boris121981@postgrespro.ru", "phone": "+70381699675"}
0077E1	2017-07-14 07:19:00-05	6000.00	0005432034890	0077E1	7650 862304	IVAN ZHURAVLEV	{"phone": "+70460520351"}
009C82	2017-07-28 19:45:00-05	6000.00	0005432035264	009C82	4287 710629	IVAN SMIRNOV	{"email": "i_smirnov-1978@postgrespro.ru", "phone": "+70345178343"}

(5 rows)

C. For the original nested loop, the best query for this should be the nest loop, because there is neither sorting nor does it fit well into memory. But the nested loop would not run for merge join even when we tried to force it to choose that join method. We tried to change the computer that we were running it on and changed the dataset to be in the same exports file and tried to do a fresh install of the lab. But none of these allowed us to run our nestloop query as a merge join. But between the two joins that ran, nested loop and hash join, nested loop performed better. We went to office hours and the TA said that this would be fine for the submission.

For the original merge join, the best query is merge join since it has the shortest runtime of 640.1 ms compared to nest loop which was 1062.6 ms and hash join which was 1024.3 ms.

For the original hash join, the best query is hash join since it has the shortest runtime of 0.464 ms compared to merge join which took 1.420 ms and nest loop which took 69.150 ms.

D.

For the original nested loop, we think that the nest loop join algorithm is better because we chose two relatively large unsorted tables to join together, so that hash join and merge join would not be favored, so the only result would be a simple scan and nest loop would be favored.

For the original merge join, we think that the merge join algorithm is the best since we made our query sort each table before running the overall join, resulting in the query favoring merge join.

For the original hash join, we think that the hash join algorithm is the best since we made our query limit each table before joining, this way the tables would fit within memory and hash join would be favored.

E.

For the original nested loop, it does match the original query optimizer that was run in step 1.

For the original merge join, it does match the original query optimizer that was run in step 1.

For the original hash join, it does match the original query optimizer that was run in step 1.

F.

For the original nested loop, we think that the original postgresql took the right decision for this query. Since our query is neither sorted nor small enough to fit in memory.

For the original merge join, we think that the original postgresql took the right decision for this query. Since our query is sorted, merge join should be favored.

For the original hash join, we think that the original postgresql took the right decision. We can say this because when Q3 (the originally hash joined query) was run using the other types of joins (the nested loop and merge joins) the query had a longer run time and was fastest when it used a hash join.